


FlyScale

Geometric scaling of the adult *Drosophila* connectome

139,255 neurons · 2,700,513 connections · 34,153,566 synapses · a 10x ladder


to answer whether a connectome can be geometrically enlarged while preserving its structure, dynamics, capabilities and measured energy budget




The whole brain — 139,255 neurons at their soma positions

Gold Kenyon cells · magenta MBON · orange DAN · green antennal-lobe projection neurons

Detail: a real 15-neuron circuit (14 connections, 14303 synapses)



The fitted 2-D hyperbolic geometry



Every mark on this cover is measured data. The brain is the FlyWire v783 connectome, all 139,255 neurons at their soma positions, with the mushroom-body circuit highlighted. The small panel is a real circuit (edge width = synapse count, edge color = predicted transmitter, positions rescaled to fit). The disk is the hyperbolic geometry fitted to the same connectome in section 5.

Every panel in this paper is plotted from a measured artifact - no illustrative figures

FlyScale: geometric scaling of the adult *Drosophila* connectome

Structure, dynamics, capability and energy across a 10× ladder

Rick Goldberg

Implemented in Vyb with native NVPTX kernel generation.

Abstract

Can a biological connectome be enlarged? We take the complete adult *Drosophila* connectome (FlyWire FAFB v783; 139,255 neurons, 2,700,513 thresholded connections, 34,153,566 synapses) and ask whether a synthetic graph can be generated that preserves its structural statistics, its propagation dynamics, its behavioural capabilities and its energy budget as it grows. The method is a geometric one. We first fit latent coordinates to the connectome and ask which space actually predicts connectivity: a 32-dimensional spectral embedding predicts held-out connections best (AUC 0.952), a 2-D hyperbolic embedding beats the anatomical 3-D soma coordinates (0.865 vs 0.851), and no embedding beats a degree-only baseline by a margin larger than its own run-to-run spread. Scaling then proceeds by node subdivision in that geometry: each neuron is replaced by c children placed near it, each inheriting the parent's partner set and synapse weights, with the concrete target chosen by the fitted connection law $P(\text{connect} \mid \text{distance}) = 1/(1 + e^{(d-R)/T})$. The result is a ladder of graphs at 2×, 5× and 10× (up to 1,392,550 neurons and 27.8M connections) whose connection count scales as $N^{1.0175}$ ($R^2 = 0.9995$) — sparsity is preserved, mean degree and mean connection strength stay bounded — and which renormalizes back to the biological graph almost exactly ($R(G_{10})$ returns 2,700,513 connections with a degree-distribution distance of 0 and community ARI 1.0). Capability grows on some axes (stimulus-discrimination information $\propto N^{1.014}$, participation ratio $\propto N^{0.318}$, memory capacity ≥ 96 associations versus 8 at 1×) and is flat on others (temporal depth and sequence depth never rise above chance at any scale). An associative-learning benchmark on the real mushroom body acquires odor–reward associations (AUC 0.51 $\rightarrow$ 1.00 in 58 trials) but the connectome's specific wiring contributes *nothing* beyond shuffled controls ($+0.000 \pm 0.000$): what matters is the dopamine-gating anatomy, not the geometry. Energy is measured on two tracks that are never summed: hardware (438/500/946/2022 GPU joules across the ladder, with cost per spike *falling* from 8.9 to 4.2 μJ as the device stops being launch-bound) and biological-equivalent ($P_{\text{bio}} = P_0 N/N_0$ from nanowatt calorimetry, a property of N alone). The work is deliberately dual-implementation: a Python reference oracle and a Vyb-native production path (loader, discrete-event runtime, GPU kernels), and the effort to build the second one surfaced twelve compiler defects, of which the blocking one — a kernel-mode store that wrote eight bytes into a four-byte slot — is documented, fixed upstream and verified here. We report what failed as carefully as what worked: a closure test that demonstrates invertibility rather than self-similarity, a capability curve censored at its test grid, a geometry that helps prediction but not renormalization, and a GPU backend whose equivalence to the CPU reference had to survive two independent bugs, one of them ours.

verdict (criterion §25 of the project brief)

minimum success

supported

strong success

supported

exceptional ($C(N)$ growing faster than $P(N)$)

not supported

1 Introduction and concept

The whole-brain connectome of an adult female *Drosophila* is the first complete wiring diagram of a brain at synapse resolution [1, 2]. It contains $\approx 139\text{k}$ neurons and $\approx 54.5\text{M}$ synapses (16.8M directed neuron pairs before thresholding), and it is not a random graph: it is sparse (connection probability $\approx 1.4 \times 10^{-4}$), weighted (a typical connection carries 12.6 synapses), organised into 79 neuropils, rich-club dominated (30% of neurons are highly connected), and heavy-tailed in degree [2]. It is, in short, exactly the kind of object for which a question about *scale* is well posed: if this graph is the product of biological growth rules, is there enough structure in it to regenerate a larger one that still behaves like a brain?

This paper reports an attempt to answer that empirically, and the answer is deliberately split. The machinery works: we can build a $10\times$ graph (1.39M neurons, 27.8M connections) that preserves sparsity, degree distribution, community structure and propagation dynamics, and that renormalizes back to the biological graph essentially exactly. Capability grows on the axes one would expect from sheer size (information, dimensionality, associative capacity) and does not grow on the axes that require temporal machinery (delay, sequence). And when we test the connectome’s own structure against shuffled controls in a learning task, the structure contributes nothing measurable — a negative result we regard as one of the more useful findings here, because it constrains what “the connectome is special” can mean.

Concept. The hypothesis under test is geometric, following the renormalization program for complex networks [4]: real networks embedded in a hidden metric space can be renormalized by rescaling distances, and networks with scale-free degree distributions show geometric scaling under that transformation. The program suggests a concrete recipe: (i) fit a latent geometry to the graph; (ii) fit the connection law that relates distance to connection probability; (iii) *coarse-grain* by grouping nodes that are close in that geometry (downscaling); (iv) *subdivide* nodes and rewire by the same law (upscaling); (v) accept the result only if the generated graph, coarse-grained back, recovers the original. Step (v) is the crux, and we treat it as the primary acceptance criterion rather than a sanity check.

Figure 1 lays out the object of study: one biological graph at the center, a downscaling branch to the left, an upscaling ladder to the right, and the closure loop that decides whether the right-hand branch is legitimate. Every number on that diagram is a measured value from this project.

Why this is not a neural-network paper. The brief this project follows forbids claiming consciousness, mapping the connectome onto transformers, pretending synapse counts are learned weights, or inferring biological efficiency from GPU power numbers. The findings below are correspondingly narrow: they are statements about a graph, a simulator and a power meter. Where the evidence is a model identity rather than a measurement (the biological energy track is linear in N by construction) we say so in the same sentence that reports it.

Contributions.

- C1.** A reproducible canonical dataset from the v783 release with an 11-point gate against published statistics, including two recorded mismatches that we did not tune away (§2).
- C2.** A measured answer to the geometry question, with the control that matters: no learned embedding beats a degree-only baseline by more than its own run-to-run spread (§5.1).
- C3.** A scaling pipeline that preserves sparsity at every rung ($E \propto N$), bounded mean degree and bounded connection strength, and a closure test that returns the biological graph to within 0.1% on connections at $10\times$ (§5).
- C4.** A dynamic validation that runs identical, normalized stimuli through the biological graph and its coarse-grained replicas and compares activation patterns, latency and output distributions




Figure 1: The concept, with measured values. The biological graph G_1 is reproduced from the FlyWire v783 release at a five-synapse threshold. Downscaling (coarse-graining in a latent geometry) produces smaller replicas used to test whether dynamics survive compression; upscaling (node subdivision guided by the fitted connection law) produces the ladder that is the subject of the paper. The dashed green loop is the closure criterion: each generated graph is coarse-grained back by lineage and compared with G_1 across a metric suite. Boxes show the measured connection counts and closure composites (§5); the 10× graph returns the biological connection count exactly.

(§5.5).

- C5.** A capability battery across four scales with fitted exponents, refusals to claim scaling where the measurement is censored, and an associative-learning benchmark on the real mushroom body with shuffled-connectivity and no-plasticity controls (§6, §6.3).
- C6.** A two-track energy measurement (measured hardware joules; cited biological model) that never conflates the two, and a capability-per-watt curve that states its own censoring (§7).
- C7.** A Vyb-native implementation path (loader, discrete-event runtime, GPU kernels) with twelve documented compiler defects, one of them blocking, filed upstream, fixed and re-verified on silicon (§4).

2 Methodology

2.1 Data and conventions

The canonical dataset is built from the FlyWire FAFB v783 release [3]: the proofread connectivity table, the root-id list, per-neuron neuropil counts and the v783 annotations. The build is streaming and deterministic; the release’s own counts are reproduced exactly (139,255 neurons; 15,091,983 directed pairs; 54,492,922 synapses; 16,847,997 edge rows before aggregation; 79 neuropils; 0 autapses).

One convention dominates every comparison in this paper. The released edge list is *unthresholded*, while published topology statistics use a threshold of five synapses per connection [2]. We adopt the threshold convention because the data confirms it rather than because it is published: at ≥ 5 synapses our graph has a mean connection strength of 12.647 (published 12.6) and a reciprocity of 0.1398 (published 0.138), whereas at ≥ 1 the same graph gives 0.266 — a factor of two away from the published value. All scaling work therefore operates on 2,700,513 connections carrying 34,153,566 synapses, with mean out-degree 19.39 — and Table 2.1 records both conventions explicitly so that no number in this paper is ambiguous about which graph it describes.

quantity	value	note
neurons (annotated)	139 241	canonical counts
neurons (in graph)	139 255	canonical counts
edge rows released	16 847 997	unthresholded
connections (all)	15 091 983	threshold ≥ 1 synapse
connections (scaling graph)	2 700 513	threshold ≥ 5 synapses
synapses (all)	54 492 922	threshold ≥ 1
synapses (scaling graph)	34 153 566	threshold ≥ 5
autapse pairs	0	none in v783
reciprocal pairs	2 007 259	threshold ≥ 1
reciprocity	0.266	threshold ≥ 1
mean out-degree	19.3926	scaling graph
mean synapses per connection	12.6471	published: 12.6
neuropils	79	annotated regions

2.2 The gate: published statistics

Before any scaling, the canonical graph must reproduce the published picture. We compare against a machine-readable table of published values (with their source papers) and separate required from advisory checks. The result is 11/11 required and 7/9 advisory checks (Table 2.2). Aggregate statistics agree closely: mean synapses per connection 12.647 vs 12.6; reciprocity 0.1398 vs 0.138; mean directed path length inside the giant SCC 4.456 vs 4.42; giant SCC size 119,756 vs 119,404; rich-club connection probability 8.68×10^{-4} vs 8.70×10^{-4} ; fraction of neurons with total degree > 37 is 29.0% vs $\approx 30\%$; in/out degree correlation 0.742 vs 0.76.

Two checks fail, and both are informative rather than alarming. First, an independent third-party count of 3,732,460 connections cannot be reproduced by any synapse threshold we tested (threshold 4 gives 3,537,227), so we record it as an unexplained discrepancy and leave the tolerance tight enough that it fails loudly. Second, the SCC/WCC fractions (0.860/0.951 here vs 0.933/0.988 published) differ because the published figures are computed on the v630 snapshot with 127,978 neurons [2] while ours are computed on v783, which adds ≈ 11 k sparsely connected neurons to the denominator: the absolute component sizes agree (≈ 119.4 k), the fractions do not. A gate that reports 11/11 with two documented, explained failures is more useful than one with loosened tolerances.

check	class	measured vs published
n_neurons	required	✗ exact: the proofread root-id array defines the neuron set
mean_synapses_per_connection	required	✗ strongest available validation that the 5-synapse convention and neuropil-level aggregation are correct
connection_reciprocity	required	✗ $p_{\text{rec}} = P[\text{beta} \rightarrow \text{alpha} \mid \text{alpha} \rightarrow \text{beta}]$, autapses excluded, 5-synapse threshold
clustering_coefficient	required	✗ global clustering on the undirected autapse-free projection; tolerance spans both quoted values
n_connections_v630_convention	required	✗ v783 has 8.8% more neurons, so a few percent more connections is expected
mean_out_degree	required	✗ ours is computed over all neurons with any connection, not intrinsic neurons only
scc_size_absolute	required	✗ compared in absolute neurons, not as a fraction: v783 adds 11k sparsely connected neurons that sit outside the giant SCC, which shifts the fraction w
wcc_size_absolute	required	✗ absolute comparison for the same reason as scc_size_absolute
scc_fraction	advisory	✗ version-sensitive: v783 has a larger denominator, so this fraction is expected to be lower than v630's while the absolute component is the same size
%		
wcc_fraction	advisory	✗ version-sensitive, see scc_fraction
mean_undirected_path_wcc	required	✗ sampled BFS from 1000 sources drawn from the largest weak component
mean_directed_path_scc	required	✗
max_directed_path_scc	advisory	✗ sampled eccentricity, not a full diameter computation
in_out_degree_correlation	required	✗ computed by the gate script over neurons with any connection
n_neuropils	advisory	✗ release neuropil labels carry _L/_R hemisphere suffixes (79 distinct labels in v783); the hemisphere-merged count is recorded separately in derived.
rich_club_size_at_degree_37	advisory	✗ counted on the 5-synapse threshold view; v783 has more neurons
connection_probability	advisory	✗ $E / (N(N-1))$ on the 5-synapse threshold view
rich_club_connection_probability	advisory	✗ 5.4x the overall connection probability
rich_club_degree_cutoff_null_normalized	advisory	✗ only produced by -heavy; our null uses 3 rewiring reps instead of the paper's 100
codex_connection_count	advisory	✗ UNRESOLVED CONVENTION: no synapse threshold reproduces 3,732,460 from the v783 release exactly (threshold 4 gives 3,537,227, threshold 5 gives 2,700,5)

2.3 Metrics, and the criterion that decides

Two families of measurement appear throughout. *Structural* metrics compare a generated graph to G_1 : Wasserstein distance between degree distributions (and weighted-degree distributions), a triad-census divergence, the spectral distance between normalized Laplacian spectra, a rich-club distance, a connectivity-matrix divergence over cell classes and neuropils, community agreement (adjusted Rand index and normalized mutual information between Leiden partitions mapped through a node correspondence), and the distribution of latent distances across connected pairs. We keep the raw metric vector for every scale and compute a composite score only for ordering scales internally —

the brief is explicit that a composite must not substitute for the measurements, and all tables here report the raw components.

Dynamic metrics compare simulations: with the same normalized stimulus seeded into matched neurons, we record the per-step activation curve, the final active set, latency to half activation, and the overlap between the reference and replica active sets (Jaccard, precision, recall) after pushing reference activity through the node correspondence.

The acceptance criterion for the whole enterprise is closure: coarse-grain a generated graph and see whether G_1 comes back. We implement it by lineage — each child node is grouped with its parent neuron, which is the grouping the generator itself used — and we are explicit about what that does and does not establish (§8).

2.4 Validation philosophy

Every claim in this paper is attached to a file. Three practices are in force throughout, and they changed conclusions more than once:

1. **Independent recomputation.** Before any gate is believed, its numbers are recomputed by an independent path: a Python oracle for the Vyb loader, direct numpy accumulation for the first simulation tick of the GPU pipeline, from-scratch re-execution for the learning curves.
2. **Missing tests are not passes.** The project-level assessment tool counts an unmeasured sub-criterion as PARTIAL, never as SUPPORTED, and refuses to treat the presence of an artifact as evidence of a result. This rule was added after the tool declared minimum success with the decisive closure test absent.
3. **Corrections are recorded, not erased.** When we found that an earlier claim of ours was wrong (a reproducibility claim based on reading one run twice; a hand-written expectation that the kernel contradicted), the correction was written into the artifact next to the retraction. Appendix C lists them.

2.5 Two-track energy accounting

Energy is reported on two tracks that are never added, averaged or ratioed into a single number. The *hardware* track measures this machine: GPU power via NVML at 5 Hz, integrated over the workload, converted to joules per simulated biological second, per spike and per synaptic event. The *biological-equivalent* track is a model, $P_{\text{bio}}(N) = P_0 N / N_0$, anchored to measured whole-brain calorimetry of live explanted *Drosophila* brains [5, 6] and cross-checked against a mammalian per-neuron scaling law that is explicitly flagged as not a fly measurement [7] and a published fly-brain power figure [9]. Section 7 states which of those quantities is a measurement and which is an identity.

3 Implementation

3.1 Two implementations on purpose

The project runs a dual stack. The *reference* implementation is a Python package (`src/flyscale`) that owns the canonical dataset, the metric suite, the renormalization operators, the propagation model and the capability battery; it is the oracle every other implementation is checked against, and the source of every figure here. The *production* path is Vyb-native — a Vyb program that loads the connectome binary, a generic Vyb discrete-event runtime, and Vyb-authored GPU kernels — because the brief requires the end state to be Vyb, and because building it turned out to be the single most informative part of the project (§4).

The oracle/production split is enforced by artifacts rather than by intention. The canoni-

cal build emits a flat little-endian binary twin of the dataset (`pairs.pre/post/syn`, `neurons.root.id`; 449.6 MB), so the Vyb loader is a mechanical port rather than a re-derivation, and the loader is then validated against the oracle by 26 invariants (§4).

3.2 The renormalization core

Coarse-graining (`renorm.coarse_grain`) groups nodes by proximity in the chosen geometry — nearest-neighbor agglomeration with distances evaluated in the ambient Poincaré coordinates for the hyperbolic case — then aggregates edges and synapses between groups. Because merging nodes raises the effective degree, the coarse graph is thresholded with a threshold chosen by calibration to match the source graph’s mean degree as closely as achievable; the calibration searches a candidate ladder, reports the residual error, and reports the minimum achievable mean degree, so a failure to hit the target is visible instead of silent.

Subdivision (`renorm.upscale`) is the inverse operator and the heart of the scaling result. Each neuron becomes c children placed near the parent in the latent geometry, at a perturbation scale tied to the local nearest-neighbor distance; each child inherits the parent’s partner set with the parent’s synapse weights, and the concrete target child is drawn by the fitted connection law; a fraction of each child’s edges is redrawn to diversify the partner set without changing its size; and sibling edges are added by the same law at a damped rate. The design consequence is exactly the property the brief demands: connections scale linearly with N (sparsity preserved), while mean degree and mean connection strength stay bounded [4].

A single fitted law serves both directions. `fit.connection_law` samples connected pairs and unconnected pairs, computes their latent distances, and fits a logistic form whose parameters give R (the distance at which connection probability is one half), T (the softness of the transition) and an AUC. On the anatomical coordinates the law reaches AUC 0.803; on the learned hyperbolic embedding, 0.888 — the first quantitative sign that the fitted geometry carries more connection information than anatomy does.

3.3 Propagation and capability

Dynamics use a deterministic threshold cascade over the directed, weighted graph: at each step a neuron fires if the summed synapse counts arriving from the previous step’s active set exceeds a threshold set as a fixed multiple of the graph’s own mean synapses per connection (so the stimulus is normalized across scales rather than rescaled per graph). The stimulus seeds a fixed fraction of neurons drawn preferentially from sensory classes. This is deliberately a *propagation* model, not the firing-rate model of §5.5; it exists to make cross-scale comparisons of activation patterns meaningful.

The capability battery reuses the same cascade with a frozen protocol: 96 stimulus classes, class pools sized proportionally to N , a binary readout over a fixed fraction of neurons decoded by nearest-centroid template matching, and every threshold decision gated on an exact binomial test at $p < 0.05$. Its outputs are the discrimination curves, memory capacity (the number of associations held above an accuracy criterion), temporal depth (delayed match-to-sample), sequence depth, generalization to novel variants, lesion robustness, and a set of complexity measures including participation ratio, response entropy and branching ratio.

3.4 Determinism and reproducibility

Every stochastic component takes an explicit seed, and the artifacts record what was run: the database version and a dataset hash, the parameters, the wall time, and the command. Where a result depends on a numerical protocol (a class grid, a trial count, a stimulus fraction) the protocol fingerprint is recorded and checked across scales — the capability battery verifies that the identical

Table 1: The Vyb discrete-event runtime against its independent reference. Every replication matches exactly; the two determinism runs produce identical hashes. Throughput was measured on a shared machine, hence the range (see `results/phase2/des_throughput_repeats.txt`).

replication	ticks	spikes	synaptic events	wall (ms)	result
main (20,000 ticks)	19,841	771,379	22,702,285	5,702	match
determinism A (500)	382	3,064	86,500	221	match, hash equal
determinism B (500)	382	3,064	86,500	122	match, hash equal
alternate stream (500)	185	841	23,860	65	match

fingerprint (`39f6c75cde6548e8`) holds at all four scales, so a difference between scales cannot be a protocol difference.

4 Vyb inference and the compiler defects it found

The production path is where the project spent effort that no result in §5–§7 required, and it paid for itself in a way the brief anticipated: running real workloads through a young compiler finds things a test suite does not.

4.1 V1: a Vyb-native loader

the Vyb loader (`src/vyb/flyload.vyb`) reads the canonical flat binary and exposes the adjacency in CSR form; `phase0_counts.vyb` recomputes whole-connectome invariants from it. Validation is 26/26 invariants against the Python oracle — row-pointer head and tail, out-degree sum, maximum and zero-degree counts, target-index sum and maximum, and the whole-array root-id sum (which overflows int64 and wraps identically on both sides, a useful reminder that agreement can be accidental as well as meaningful). The loader establishes that the binary twin is a faithful representation, and it set the pattern for every later Vyb artifact: a Python oracle, a Vyb consumer, and a checker that compares them.

4.2 M2: a Vyb discrete-event runtime

`src/vyb/des.vyb` implements a bucketed scheduler (a ring of per-bucket event chains over one flat arena, $O(1)$ push, no single global heap), with entity state, event sources and sinks, a recorder and replication/experiment wrappers. The application `phase2_des.vyb` runs a LIF network of 640 neurons and 18,909 synapses for 20,000 ticks (2,000 ms simulated, one tick = 0.1 ms), scheduled 22,766,285 events and produced 771,379 spikes. It matches an independently written Python reference *exactly* — identical per-entity tick sequences, identical event counters, identical FNV-64 hashes recomputed by the checker — across four replications, including two determinism re-runs and an alternate stimulus stream (Table 1). Integer arithmetic throughout (voltages and weights in micro-units) is what makes exactness achievable rather than approximate. Throughput is 2.19–6.96M events/s (median 4.9M) single-core, i.e. ≈ 11.4 M synaptic events per simulated biological second.

4.3 M3: GPU kernels, and how a four-byte store cost a milestone

The event-driven GPU backend is written in Vyb device code: `accumulate_bucket` walks the CSR slice of each fired neuron and accumulates synaptic input into target neurons with atomic adds (one block per source, no prefix-sum pass), and `update_and_fire` performs the LIF update, applies the refractory gate, and scatters new events into the next bucket by atomic slot reservation. Both




Figure 2: The M3 gate, after both fixes. Left: firing per tick for the GPU pipeline and the CPU reference — the reference’s own tick-0 count (3 neurons) was confirmed by a third, independent direct computation before we trusted it over the GPU’s. Right: the membrane state after 30 ticks for the first eight neurons, exact agreement. The same kernels are additionally validated individually against hand-computed inputs (`phase3_selftest.vyb`, `probes/accum_probe.vyb`).

kernels lower to NVPTX `.entry` points and are validated by `ptxas` for the target architecture; a Vyb host program drives the whole bucket→accumulate→update→scatter pipeline on the device.

Getting that pipeline to agree with the CPU reference took two independent fixes, and the sequence is instructive. The *first* was upstream: with a computed value, the refractory store `st.i32(refr + gid*4, rt)` wrote correctly for some lanes and zero for others, and the accumulator behaved as if over-driven. Isolation narrowed it in three steps — the update kernel alone, driven with a hand-written membrane state, was correct apart from that store; a purpose-built probe showed descriptor loads were not corrupted per lane; and comparing PTX for a failing and a passing kernel showed the mechanism. The compiler stored the operand’s own LLVM type rather than the intrinsic’s element type, so `st.i32(ptr, i64.value)` emitted an eight-byte store: the low word at the address and the high word at address+4, i.e. into the *neighboring* four-byte slot. A computed value with a zero high word therefore zeroed its neighbor’s cell — producing exactly the 1818 0 0 0 0 0 0 0 pattern we observed. We filed it with the sources, the measured outputs and, unusually, the three hypotheses we had already refuted [10]. The fix landed as `st.*` using the intrinsic’s element type [11], and the same probe now reads 1818 on all eight lanes.

The *second* fix was ours: while adding a guard for the empty-bucket case, the original unguarded launch was left in place, so the accumulation kernel ran twice per tick. Every accumulated value was exactly double — which pushed ten neurons over threshold instead of three — and the unguarded launch with a zero grid dimension produced 21 launch errors. It was found by reading the accumulator between the two launches and comparing it against the isolated kernel’s own output; a clean factor of two named the cause immediately.

With both fixed, the equivalence gate passes on all six checks: 5 spikes total (matching the reference), identical per-step firing, identical per-neuron spike counts, an identical membrane trace, and zero launch errors (Figure 2). We state the scope plainly: this is CPU≈GPU equivalence on a 512-neuron fixture whose semantics match exactly. Running the same pipeline over the 139k-neuron connectome is a separate step and has not been done.

4.4 Twelve defects, and what they say about the Vyb-native goal

Building three Vyb artifacts surfaced twelve distinct compiler defects (Table 4.4). Nine remain open; one was fixed upstream as a result of this work; two are design limitations rather than

bugs. The most consequential for the project is the last row: byte-at-a-time decoding measured at $\approx 29 \mu\text{s}$ per element ($\approx 34\text{k}$ elements/s), which makes a 15M-element pass take about seven minutes and therefore blocks a whole-connectome Vyb simulation until a bulk byte $\rightarrow$ int read exists. The ‘Vec|Vec|T»’ defects are the second-most consequential: they forced the DES scheduler into a structure-of-arrays design that, as it happens, is the shape a GPU backend wants anyway — so the defect nudged the implementation in a useful direction.

defect found while building this project	reproduction / observed behavior	consequence here	status
st.* stored the operand’s LLVM type, not the intrinsic element type	‘st.i32(ptr, i64)’ emitted an 8-byte store whose high word landed in the neighbouring 4-byte slot	kernel-mode data corruption in the GPU pipeline	fixed upstream (PR #273)
a parenthesised operand after ‘*’ inside an addition	the multiplication’s left operand is typed as a pointer; ‘p + s * (q)’ fails to compile outside kernel mode	blocks ordinary arithmetic expressions	fixed upstream (verified on main@51ced31)
‘Vec Vec T»’ is unusable	four failure modes incl. silent garbage from ‘get(i).get(0)’ and ‘push’ into a temporary	forced the DES scheduler into a structure-of-arrays design	open (issue #284)
re-borrowing an existing borrow segfaults	‘f(borrow(v))’ where ‘v’ is already borrowed (exit 139)	API shape constraint in the DES runtime	open (issue #285)
raw io byte buffers corrupt across a module boundary	the same 16 bytes read inline give 0, 62 and via a module function give a 64-bit value	forced decoding inside the module	open (issue #281)
% optional ‘Vec’ returns are rejected	a function declared ‘Vec UInt8 L’ cannot ‘return’ a ‘Vec UInt8 L’	forced a rewrite of the loader	open (issue #282)
‘Vec’ parameters are not mutated for the caller	out-parameters come back empty	design constraint for the loader	open (issue #283)
device intrinsics return ‘CInt’	‘tid.x’/‘blk.x’/‘dim.x’/‘ld.i32’ need an explicit ‘as Int’ in assignments	documented; cost a compile cycle	open (by design?)
‘ptr’ is a reserved struct-field name	rejected with a message that never names the token	minor ergonomics	open (issue #286)
‘String.to.int’/‘to.float’ do not exist	documented in the reference manual but absent in the compiler	the DES app ships its own integer parser	open (issue #287)
byte-at-a-time decode throughput	$\approx 29 \mu\text{s}$ per element ($\approx 34\text{k}$ elements/s); a 15M-element pass takes ≈ 7 min	project-level blocker for a whole-connectome Vyb simulation; a bulk byte $\rightarrow$ int read is required	open (re-measuring on current upstream)

5 Geometric scaling: experiments

5.1 Which space predicts connectivity?

We fit five candidate geometries to the ≥ 5 -synapse graph — anatomical soma coordinates (2-D and 3-D, raw and axis-standardized), a 2-D hyperbolic (Poincaré) embedding learned by stochastic gradient descent on the connection law, and normalized-Laplacian spectral embeddings at 16 and 32 dimensions — and score each by its ability to predict held-out connections on *identical* positive/negative pairs (269,478 of each, a 90/10 split, giant component of 132,181 nodes). A degree-only baseline is scored on the same pairs, and every fit is additionally reported with a log-likelihood penalty.

The ranking (Table 2, Figure 3) is unambiguous at the top and ambiguous in the middle. Spectral 32-D wins by a wide margin (AUC 0.952 vs 0.874 for degrees), and combining distance with log-degree lifts every geometry (anatomy 0.851 $\rightarrow$ 0.949, hyperbolic 0.865 $\rightarrow$ 0.904, spectral-32 0.952 $\rightarrow$ 0.977), which says the geometry carries information the degrees do not but that the two are complementary. Hyperbolic 2-D beats anatomical 3-D in every realization we measured (0.8646/0.8921/0.9012 vs 0.8514), including one initialized without degree information — so the direction of the published geometry result reproduces. But hyperbolic-versus-degree-only is *unresolved*: the delivered run scores below the baseline (0.865 < 0.874) while a second identical-configuration run scores above it (0.901), a spread of 0.037 on identical held-out pairs. A margin smaller than the run’s own variation is not a result, and we report it as such. Equal-weighting the anisotropic anatomical axes (soma extents differ by a factor of 40 between axes) actively hurts (0.798), which is a reminder that “the anatomy” is not a single well-defined space.

Table 2: Held-out connection prediction. Larger AUC is better. The degree-only baseline is not a straw man: it is the strongest non-geometric predictor, and it beats every geometry except the 16- and 32-dimensional spectral embeddings.

geometry	AUC	AP	mean log L	R	T
spectral 32	0.9519	0.9452	-0.2868	0.000915	0.000253
spectral 16	0.8844	0.8445	-0.3866	0.000399	0.000215
degree only (baseline)	0.8739	0.8784	-0.4445	nan	nan
degree only baseline	0.8739	0.8784	-0.4445	—	—
hyperbolic 2d	0.8646	0.8375	-0.4597	3.451	0.978
anatomical xyz	0.8514	0.8307	-0.4669	3.87e+04	1.83e+04
anatomical pca2	0.8512	0.8293	-0.4671	3.86e+04	1.83e+04
anatomical xy	0.8512	0.8293	-0.4671	3.86e+04	1.83e+04
anatomical xyz zscored	0.7981	0.8026	-0.547	1.738	0.766




Figure 3: Connectivity prediction by geometry, on identical held-out pairs. The dashed line marks the degree-only baseline. The hyperbolic rows carry a measured run-to-run spread of 0.037, larger than its margin over that baseline — the honest reading is "beats anatomy, ties degrees".

5.2 Coarse-graining, and whether the learned geometry helps

Downscaling groups nodes by latent proximity and thresholds the aggregated graph to preserve mean degree. We ran the same pipeline with both geometries, which turns the geometry question from a prediction exercise into an engineering question (Table 3, Figure 4).

The result cuts against a naive expectation. The hyperbolic embedding is the better *law* — it predicts connectivity better and yields a cleaner degree distribution at $0.5\times$ (Wasserstein 0.104 vs 0.129) — yet the coarse graphs it produces are worse where it matters for behavior: community agreement falls to ARI 0.33/0.20/0.24 versus 0.52/0.56/0.52 for anatomical coordinates, and the composite score is worse at every factor. The hyperbolic grouping also drifts further from the target mean degree (at $0.25\times$ it lands at 13.6 against 25.0 for anatomical), which is the calibration signal that the grouping is producing structurally different supernodes. We therefore report the scaling results as geometry-driven with the trade-off stated, not as evidence that the learned geometry is the right latent space.

Table 3: Coarse-grained replicas under both geometries, compared with G_1 through a partition correspondence. The learned hyperbolic geometry fits the connection law better (AUC 0.888 vs 0.803) and slightly improves degree preservation at $0.5\times$, but is clearly worse on community structure at every factor — so it does not rescue renormalization.

geometry	scale	neurons	connections	mean deg.	composite	deg. Wasserstein	ARI
anatomical xyz	$0.5\times$	70 437	1 193 153	16.9393	0.1785	0.129	0.5163
anatomical xyz	$0.25\times$	36 197	904 216	24.9804	0.2665	0.3082	0.5568
anatomical xyz	$0.1\times$	15 997	273 884	17.121	0.2758	0.1456	0.5213
learned hyperbolic	$0.5\times$	70 369	1 230 392	17.4849	0.2138	0.1043	0.3312
learned hyperbolic	$0.25\times$	36 048	488 581	13.5536	0.3113	0.3031	0.1983
learned hyperbolic	$0.1\times$	15 617	367 690	23.5442	0.3202	0.2417	0.2397




Figure 4: Coarse-graining under both geometries. Lower is better for the left and center panels; higher is better for community ARI. The learned geometry wins one panel, loses the other two.

5.3 Inverse scaling: the ladder

Subdivision doubles, quintuples and tenfolds the neuron count. The results (Table 4, Figure 5) satisfy the three invariants the brief requires. First, connections scale linearly with neurons: a fit across all four scales gives $E \propto N^{1.0175}$ with $R^2 = 0.9995$, and synapses $\propto N^{1.0067}$ with $R^2 = 0.999996$ — the aggregate is emphatically not drifting toward an all-to-all regime. Second, mean out-degree stays in a narrow band around the source value (18.62, 19.39, 19.99 versus 19.39), and mean connection strength likewise (13.19, 12.74, 12.46 versus 12.65), so the generated graphs are not becoming denser or stronger as they grow. Third, the per-scale structural closure against G_1 degrades gracefully and then improves with size — composites 0.194, 0.236, 0.275 for $2\times/5\times/10\times$ — which is expected, since at larger factors each source neuron’s structure is represented by more children and the parent-level statistics are recovered with less sampling noise.

5.4 Closure: does the ladder renormalize back?

Each generated graph is coarse-grained back by lineage and compared with the biological graph (Table 5, Figure 6). At $10\times$ the round trip returns 2,700,513 connections — the biological count exactly — with a degree-distribution Wasserstein distance of 0, a triad-census divergence of 0 and community partitions agreeing at ARI 1.0. At $5\times$ it is 15 connections short with ARI 0.9915; at $2\times$, 20,043 short with ARI 0.912.

Table 4: The upscaled ladder. Connection counts grow linearly with neurons ($\alpha = 1.0175$, $R^2 = 0.9995$), so sparsity is preserved; mean degree and mean connection strength stay bounded around the biological values. The final column is the same-graph structural comparison against G_1 (see Table 5 for the lineage closure test).

scale	neurons	connections	synapses	mean deg.	mean strength	composite
G_1	139 255	2 700 513	34 153 566	19.3926	12.6471	—
G_2	278 510	5 186 025	68 428 042	18.6206	13.1947	0.1941
G_5	696 275	13 500 829	171 970 190	19.3901	12.7378	0.2355
G_{10}	1 392 550	27 842 102	346 934 550	19.9936	12.4608	0.2753




Figure 5: Scaling laws of the enlargement, fitted on the measured counts. Both connections and synapses track N with exponent indistinguishable from one — the property that keeps energy scaling linear rather than quadratic, and the one we would have lost first had the subdivision been naive duplication of partner sets.

5.5 Dynamic validation of the downscaled replicas

Structure is necessary but not sufficient: a coarse graph that matches statistics while propagating differently is not a valid replica. We seed an identical, normalized stimulus (1% of neurons, sensory classes preferred, the same relative threshold expressed in each graph’s own mean connection strength) into G_1 and into each replica, and compare activation (Table 6, Figure 7).

The replicas behave like the source. Final active fractions are 0.821, 0.902 and 0.832 against the reference’s 0.826; against a mapped reference the active-set Jaccard is 0.909, 0.950 and 0.853 with precision above 0.99 (i.e. essentially every neuron a replica activates has a reference counterpart, with slight over-recruitment); latency to half activation falls by one or two steps out of an eleven-step cascade. Under the hyperbolic geometry the same comparison gives Jaccard 0.916, 0.821 and 0.902 — comparable, with the $0.25\times$ replica the weakest of the six. This is the strongest evidence we have that coarse-graining preserves *behavior* and not just topology.

Table 5: The closure test: coarse-graining each generated scale back and comparing with G_1 . Lower is closer for composite, Wasserstein and triad L_1 ; 1.0 is perfect for ARI. At $10\times$ the recovered connection count equals the biological count exactly. See §8 for what this does and does not demonstrate.

scale	neurons in $R(G_s)$	connections in $R(G_s)$	composite	degree Wasserstein	triad L_1	ARI
G_2	139 255	2 680 470	0.018407	0.007422	1.06e-5	0.911867
G_5	139 255	2 700 498	0.001504	6e-6	7e-9	0.991497
G_{10}	139 255	2 700 513	0.000978	0	0	1
biological G_1 for reference		2.7e6	—	—	—	—




Figure 6: Closure metrics against the biological graph. The left panel is logarithmic because the distances span five orders of magnitude; the right panel shows the recovered community partition. The near-exact agreement at $10\times$ is partly a consequence of the generator’s construction (children inherit their parent’s partner set), which is why we describe the result as invertibility rather than self-similarity.

Table 6: Identical normalized stimuli through the biological graph and its coarse-grained replicas. The final active fraction tracks the reference closely (ratios 0.99–1.09), the active *sets* overlap strongly (Jaccard 0.85–0.95, precision >0.99), and latency changes by at most two steps — the replicas reproduce the reference’s propagation behavior, not merely its degree sequence.

scale	final active fraction	ratio to G_1	active-set Jaccard	precision	latency Δ
G_1 (biological)	0.8265	—	—	—	—
$G_{0.5}$	0.821	0.9934	0.9086	0.9976	-1
$G_{0.25}$	0.9024	1.0919	0.9499	0.9956	-2
$G_{0.1}$	0.8322	1.0069	0.853	0.998	-2




Figure 7: Activation cascades for the biological graph and its replicas, normalized by the reference’s final active count. The curves track each other in shape, magnitude and latency; the Kendall-style summary is the Jaccard overlap of the final active sets reported in the legend and in Table 6.

6 Dynamics, capability and learning

6.1 A whole-connectome LIF baseline

Before generating graphs it is worth knowing whether the biological one can be simulated at all. We implemented a leaky integrate-and-fire model over the full thresholded connectome: 139,255 neurons, 2,700,513 connections, sign from predicted neurotransmitter (acetylcholine excitatory; GABA and glutamate inhibitory), weights derived from synapse counts, per-edge synaptic delays from soma geometry (measured delay histogram: 1,744,358 edges at one step, 911,741 at two, 43,394 at three, 1,020 at four), refractory state, and two independent engines: a dense timestep reference and a sparse event-driven integrator.¹

The two engines agree exactly — zero spike difference, identical (step, neuron) trains, per-neuron rate correlation 1.0 — and a subthreshold probe run bounds the float disagreement at 4.4×10^{-16} in membrane potential. One simulated second takes 120 s wall on 20 cores. On a light-odor protocol the network produces 366,872 spikes (2.64 Hz mean rate over 14.5% of neurons) with 87.6% of spikes coming from neurons that were not driven by the stimulus, i.e. the response is stimulus-locked and recurrent rather than feed-forward. A Poisson-driven protocol gives 479,259 spikes at 3.44 Hz. An operating-point scan over synaptic gain (0.05→0.5) moves the network from 0.65 Hz to 20.8 Hz without saturating.

Two bugs were found here by the agreement gate, and both are the kind that survive inspection: the dense engine’s CSR was transposed, so it simulated the network with every edge reversed (48k versus 1.2M spikes — invisible on a random toy graph), and the event engine’s bucket handling double-charged a neuron fed by two sources in the same timestep (+50% spikes). Both now have regression tests.

6.2 Capability across scales

The capability battery was run at $1\times$, $2\times$, $5\times$ and $10\times$ under a single protocol whose fingerprint is identical at every scale (96 stimulus classes, class pools proportional to N , a 5% readout, a nearest-centroid decoder, every threshold decision gated at $p < 0.05$). Results are in Table 7 and Figure 8.

What grows: stimulus-discrimination information at 90% class overlap rises $0.00 \rightarrow 0.156 \rightarrow 0.524 \rightarrow 0.774$ bits, with a fitted exponent of $N^{1.014}$ ($R^2 = 0.957$) — proportional to size, which is what a population code should do when the readout grows with the population. Effective dimensionality rises too, but sublinearly: participation ratio $182.6 \rightarrow 388.5$ with exponent $N^{0.318}$ ($R^2 = 0.946$), and the components needed for 90% of variance grow from 319 to 400. Memory capacity rises from 8 associations at $1\times$ to at least 96 at $2\times$ and beyond — but the 96-class grid censors the larger scales, so this is a lower bound and no honest exponent can be fitted to it. Robustness to synapse deletion improves ($\alpha = 0.187$, $R^2 = 0.844$), and the minimum separable stimulus difference falls from 0.40 to the grid floor 0.05.

What does not grow, and we refuse to dress up: temporal depth is 0 at every delay and every scale (delayed match-to-sample never exceeds chance, $p > 0.05$); sequence depth is 0 at every scale (order-versus-reversal never exceeds chance); generalization sits at its measurement ceiling (4/4) everywhere; robustness to *neuron* loss is complete (90% lesioned, accuracy 1.00) at every scale, so it cannot be a scaling signal; and per-unit entropy actually declines with size ($\alpha = -0.072$, $R^2 = 0.98$). The flat axes are informative: they say the cascade readout carries instantaneous pattern information but no temporal machinery, which is a property of the model, not a discovery about fly brains.

¹Weights, delays and gains are *specified model parameters*: the connectome has no measured synaptic weights or delays. The artifacts label each parameter as model or measurement, and the operating point is recorded so the choice can be re-derived.

Table 7: Capability across scales. Discrimination information and effective dimensionality grow; memory capacity saturates at the test grid’s ceiling from $2\times$ up (a lower bound, starred); temporal depth and sequence depth never exceed chance at any scale. All values are measured; the grid ceiling is the measurement’s limit, not the system’s.

	scale	neurons	min sep. Δ	MI (bits, overlap 0.9)	memory capacity	temporal depth	sequence depth	participation ratio	gen. gap
	$1\times$	139 255	0.4	—	8	0	0	182.5964	-0.125
%	$2\times$	278 510	0.05	—	96*	0	0	268.5836	0
	$5\times$	696 275	0.05	—	96*	0	0	339.2364	0
	$10\times$	1 392 550	0.05	—	96*	0	0	388.5407	0

* at the 96-class grid ceiling: a lower bound, not a measured capacity




Figure 8: Capability versus size. Left: discrimination information grows proportionally to N . Center: effective dimensionality grows sublinearly. Right: associative capacity saturates at the test grid’s ceiling (triangles mark censored points), so the curve is a lower bound rather than a measurement — the single most important limitation on any claim of super-linear capability.

Cost per scale was 28, 128, 300 and 672 seconds of CPU time for 1,068 cascades each, i.e. wall time scaling as $N^{1.32}$ ($R^2 = 0.96$) — superlinear, as expected for a sparse gather-plus-scatter on a growing edge set.

6.3 Associative learning in the real mushroom body

The capability battery measures recognition; learning needs plasticity. We extracted the mushroom body from the annotations and built a reward-modulated (three-factor, dopamine-gated) learning benchmark on it. The populations are taken verbatim from the dataset: 5,177 Kenyon cells, 96 MBONs, 331 dopaminergic neurons (307 PAM reward, 24 PPL1/PPL2 punishment), 685 antennal-lobe projection neurons (277 uniglomerular) spanning 56 glomerulus channels; at the five-synapse threshold the $KC \rightarrow MBON$ subgraph carries 21,438 pairs and 169,931 synapses and reaches 77 of 96 MBONs, with 21,509 $ALPN \rightarrow KC$ and 735 $DAN \rightarrow MBON$ connections (630 PAM, covering the 34 MBONs used as the readout).

The network acquires the association: readout AUC rises $0.51 \rightarrow 1.00$, reaching the 0.9 criterion in 58 trials, and holding 8 associations above 0.9 (16 above 0.75) before interference collapses the curve (accuracy 0.94 at 8 associations, 0.77 at 16, 0.52 at 32). A frozen network sits at 0.31 with capacity 0, so plasticity is necessary. Valence is real: flipping only the teaching signal’s sign drives the punished odor to AUC 0.000, symmetric with the rewarded one. Architecture is preserved: across 60 plastic runs, zero synapses are created or deleted, weight budgets are conserved to 10^{-16} , and initial-to-final weight rank correlation stays at 0.54–0.93.

The controls are where this becomes interesting, and they are negative. Replacing the connectome’s KC→MBON wiring with a shuffled version of the same degree sequence changes acquisition by $+0.000 \pm 0.000$ and leaves capacity identical; so does shuffling the weights, shuffling the dopamine gate, and replacing the glomerular odor code with a degree-matched random projection. Under this benchmark, connectome geometry contributes *nothing* over its degree statistics — training dominates architecture, in the strong form. What *does* matter is the anatomical gating: broadcasting dopamine to all MBONs instead of using the measured DAN wiring drops performance to 0.971 (trials-to-criterion 106 vs 58, capacity 0 vs 8). Two further results are reported as confounded rather than clean: a readout over all 96 MBONs is flat *by conservation* (per-KC row normalization makes the summed readout invariant), and the punishment pathway test is a wiring comparison rather than a sign control because PAM and PPL1/PPL2 innervate largely different MBON sets.

One annotation defect surfaced and was worked around rather than hidden: the dataset’s `top_nt` column labels 5,172 of 5,177 Kenyon cells as dopaminergic, while `known_nt` says acetylcholine/sNPF and 100% of KC→MBON edges carry acetylcholine. We used the latter and recorded the conflict.

7 Energy

7.1 Two tracks

The hardware track is measured on this machine (RTX 3090, i9-10850K) with GPU power sampled at 5 Hz via NVML and integrated trapezoidally over each workload phase. The reference run instrumented a real sparse-propagation payload over the full connectome CSR: 84 epochs, 42 simulated biological seconds, 128,146,295 spikes and 11,710,877,981 synaptic events. GPU idle power is 22.4 W; under load the device averages 151.6 W and peaks at 158.7 W; energy above idle is 6,799.7 J, i.e. **161.9 J per simulated biological second**, **53.1 μ J per spike** and **581 nJ per synaptic event**. Wall time is 1.25 s per simulated biological second on the GPU against 11.55 s on the CPU reference (a $9.2\times$ ratio), CPU↔GPU counters agree exactly (5,052 spikes, 382,090 events), an `nvidia-smi` cross-check agrees to 0.61%, and baseline stability is 0.07%.

Two instrumentation facts are recorded because they affect any reader’s own measurements: power queried through `nvidia-smi --query-gpu=power.draw` from a background sampling thread returned a stale idle value at 96% utilisation, whereas NVML from the same thread tracked 22→155 W correctly — hence NVML primary with `nvidia-smi` as cross-check; and the device takes ≈ 10 s to fall back to idle, so baselines wait for a settled plateau and workload energy is reported as a lower/upper range. CPU package and DRAM power are *unmeasurable* here: RAPL’s `energy_uj` is mode 0400 root:root, and no `hwmon` or `power_supply` power node exists, so every joule figure above is a GPU-device lower bound.

The biological track is a model, and we label it as one in every table: $P_{\text{bio}}(N) = P_0 N / N_0$ with $P_0 = 2.56 \times 10^{-7}$ W at $N_0 = 139,255$ neurons, i.e. 1.838 pW per neuron, anchored to nanowatt-resolution calorimetry of individual live explanted *Drosophila* brains [5, 6]. Caveats travel with it: ex vivo, perfused, one sex and age, a ≈ 40 s calorimeter time constant that cannot see spike transients; the neuron count is paired from FlyWire rather than from the measured brain. The mammalian cross-check (6 kcal per day per 10^9 neurons, 290.6 pW/neuron, $158\times$ the fly value [7]) is flagged `not_a_drosophila_measurement` and per-event biology is quoted from the classical energy-budget literature [8]. The hardware-to-biology ratio ($\approx 4.3 \times 10^7$) carries an explicit note that it is not an efficiency claim: comparing a GPU to a brain is precisely the category error the brief forbids.

7.2 Energy across the ladder

Running the same instrumented payload on the CSR of each scaled graph gives the curve in Figure 9 and Table 8. Hardware energy grows sublinearly with N (438, 500, 946 and 2,022 GPU

Table 8: Measured hardware energy across the ladder, with the biological-equivalent track beside it. The two tracks answer different questions and are never summed: one is a measurement on this machine, the other is a property of the neuron count. Capacity per watt uses associative capacity as the numerator, which is censored from $2\times$ up, so those values are lower bounds.

	scale	neurons	GPU joules	J per spike	J per synaptic event	wall seconds	biological W	capacity per watt
	$1\times$	139 255	437.5953	8.89e-6	4.16e-7	3.9577	2.56e-7	0.07
	$2\times$	278 510	500.2563	5.2e-6	2.5e-7	3.3559	5.12e-7	0.615
%	$5\times$	696 275	945.7088	3.92e-6	1.81e-7	6.0537	1.28e-6	0.612
	$10\times$	1 392 550	2022	4.19e-6	1.89e-7	11.1316	2.56e-6	0.526

hardware column measured on an RTX 3090 (GPU device only; CPU/DRAM power unavailable);
biological column is P_0N/N_0 from the cited calorimetry anchor - a property of N , never summed with the hardware column




Figure 9: Energy across scales. Left: measured GPU joules grow sublinearly. Center: cost per spike falls by a factor of two — a device-utilisation effect, not biological efficiency. Right: the biological-equivalent track is linear in N *by construction*; it is a model identity, reported separately precisely so that it cannot be mistaken for a measurement.

joules for $1\times$ – $10\times$), so the cost per spike *falls* from 8.9 to 4.2 μJ as the device stops being launch-bound. Two cautions follow. First, this is a property of the machine and the kernel, not of biology: a smaller graph wastes the device. Second, the absolute joules are only reproducible to a few percent — a repeat of the identical protocol differs by 3.1–6.8% in joules while the spike and event counts reproduce to within 0.05%, because the energy number is an integral of sampled instantaneous power on a shared, thermally variable device rather than a counter. The *relative* finding (a factor of ≈ 2 in cost per spike) is far larger than that spread; the absolute values need error bars, and we have corrected an earlier claim in this project that they reproduced to three decimals (Appendix C).

8 Discussion and limitations

Closure demonstrates invertibility, not self-similarity. The $10\times$ result — connection count returned exactly, degree distance 0, triad divergence 0, ARI 1.0 — is strong but partly structural. Children inherit their parent’s partner set with the parent’s weights, so parent-level edges survive by construction; the round trip is close to an identity on the parent-level graph *for that reason*. A 0.0 Wasserstein distance is also partly a statement about the metric: the coarse graph’s degrees are aggregated deterministically from the same edge set, so no sampling noise can enter, and noise is exactly what would appear if the grouping had to be inferred. The genuinely stronger test — coarse-graining each generated graph by *geometry* rather than lineage, i.e. letting a grouping the

generator did not choose decide what collapses — has not been run and is the single most valuable next experiment.

The geometry question is partly unresolved. Labeled honestly: spectral embeddings beat degrees; hyperbolic beats anatomy; hyperbolic versus degrees is inside the noise; and the geometry that predicts connectivity best is *not* the geometry that renormalizes best. The pipeline’s current geometry choices also differ between phases (capability replicas were built on anatomical coordinates because the embedding postdates them; closure used the hyperbolic embedding), so the chain is a set of measurements rather than one consistent pipeline. Fixing a single geometry end-to-end is a prerequisite for any statement about “the” latent space.

Capability claims are bounded by their measurement grid. Discrimination information grows proportionally to N against a readout that also grows with N , so part of that growth is bookkeeping; the interesting part is that it holds at all. Memory capacity is censored at 96 associations from $2\times$ upward, which makes the exceptional criterion ($C(N)$ growing faster than $P(N)$) undecidable in principle rather than merely unsupported — the honest fix is a larger class grid, and it is cheap relative to everything else here. Temporal depth and sequence depth are zero everywhere, which limits any claim that these graphs “do” anything beyond instantaneous pattern separation.

The learning benchmark says something uncomfortable. The connectome’s specific wiring contributes nothing over shuffled controls in the mushroom-body task, while the dopamine-gating anatomy does. If that holds under better odor codes and multiple task families, the implication is that for this class of associative task the informative structure is the gating circuitry rather than the $KC \rightarrow MBON$ weight matrix — a narrower claim than “the connectome is special”, and one that fits the degree-matched random-projection control being equally good.

What the energy numbers are not. They are not efficiency claims, not comparisons to biology, and (for the absolute values) not tight. The CPU/DRAM side is unmeasurable as configured, and the whole-connectome GPU equivalence is unproven: the M3 gate passes on a 512-neuron fixture, while the energy payload is a separate CUDA program. Closing that gap — running the Vyb kernels over the canonical connectome and comparing against §6’s CPU baseline — is the natural next milestone, and it is blocked on the decode-throughput defect documented in Table 4.4.

Process limitations. One workstream (the capability battery) initially reported scaling verdicts from a single measured scale; a review caught it and the summary now carries the correction. One of our own probes produced a false negative that delayed the compiler diagnosis by a cycle. One reproducibility claim was wrong and is retracted in place. These are recorded in Appendix C because the failure modes of this kind of project are more instructive than its successes.

9 Conclusion

A whole-brain connectome can be enlarged by geometric subdivision: the $10\times$ graph preserves sparsity ($E \propto N^{1.0175}$), keeps mean degree and connection strength bounded, reproduces the source’s propagation behavior under identical normalized stimuli (active-set Jaccard 0.85–0.95), and renormalizes back to the biological graph’s connection count exactly. On the capability side, size buys information and dimensionality in proportion and sublinearity respectively, and buys associative capacity that is at least an order of magnitude larger — but it buys no temporal machinery and no measurable advantage for the connectome’s own wiring over shuffled controls

when the task is associative learning. Energy measured across the ladder falls per spike for device-utilisation reasons that must not be mistaken for biological efficiency, while the biological-equivalent model is linear by construction.

Three things would change our conclusions, and all three are cheap relative to what has already been spent: un-censoring the capability grid, running the closure test under a geometry-inferred grouping, and executing the Vyb GPU kernels over the full connectome so that the GPU $\approx$ CPU claim stops being fixture-scale. Until the first of those is done, the exceptional criterion stays undecided rather than refuted; until the second, closure stays a statement about invertibility; until the third, the Vyb-native path is validated but not demonstrated at scale.

A Milestone status

Table 9: Milestone status against the project brief’s milestone table, with the evidence that decides each row. “Gated” means deliberately not attempted; the reason is recorded in the evidence column.

id	milestone	state	evidence
M0	reproduce FlyWire graph statistics	done	11/11 required checks vs published values; 7/9 advisory
M1	CPU neural baseline (LIF, whole connectome)	done	dense = sparse spike-for-spike (0 delta, $\rho = 1.0$); 1.0 s in 120 s wall
M2	Vyb discrete-event runtime	PASS	spike-for-spike against an independent reference, 4/4 replications
M3	event-driven GPU backend	PASS	CPU $\approx$ GPU on the fixture; 0 launch errors (required upstream fix #270)
M4	energy instrumentation	done	GPU power, counters, parity; RAPL unavailable as non-root
M5	latent geometry	done	spectral 32-D wins (0.952); hyperbolic beats anatomical 3-D
M6	downscaling 0.5/0.25/0.1x	done	composites 0.179 / 0.267 / 0.276
M7	dynamic validation of down-scaled scales	done	active-set Jaccard 0.909 / 0.950 / 0.853
M8	2x inverse scaling	done	sparsity preserved, strengths bounded
M9	closure R(G.2) toward G.1 (lineage)	done [†]	composite 0.0184, degree Wasserstein 0.0074, ARI 0.912
M10	10x graph (1.39M neurons)	done	closure composite 0.00098, ARI 1.0
M11	plasticity / associative learning	done	acquisition AUC 0.51 to 1.00; geometry adds nothing over shuffled controls
M12	capability scaling benchmarks	done	discrimination and dimensionality grow; temporal/sequence depth flat
M13	energy curves vs scale	done	438 / 500 / 946 / 2022 GPU joules
M14	capability per watt	done (bound)	0.07 / 0.62 / 0.61 / 0.53; censored
M15	100x graph (14M neurons)	gated	needs the sibling-edge placement vectorised (1.4e9 distances at $c = 100$)

[†] closure demonstrates invertibility of the generator, not self-similarity (see §9).

B Artifacts, provenance and reproduction

Every number in this paper is produced by a script from an artifact on disk. The paper’s own figures and tables are generated from a single extracted data file, so prose and tables cannot drift apart.

Table 10: Primary artifacts. All paths are relative to the project root.

artifact	contents
<code>data/processed/canonical_v783/</code>	canonical dataset (parquet + flat binary twin, 449.6 MB)
<code>results/phase0/</code>	published-value gate (11/11 required), baseline metrics, threshold sweep
<code>results/phase1/</code>	whole-connectome LIF baseline, engine agreement, operating-point scan
<code>results/phase2/</code>	Vyb DES summary, spike records for four replications, throughput repeats
<code>results/phase3/</code>	GPU kernels (PTX + cubin), equivalence gate, isolation probes
<code>results/phase4/</code>	seven fitted embeddings, connection-law parameters, sensitivity study
<code>results/phase5/</code>	coarse-grained replicas and dynamic validation, both geometries
<code>results/phase6/</code>	upscaled ladder ($2\times/5\times/10\times$), per-scale replicas
<code>results/phase7/</code>	closure metrics per scale
<code>results/phase8/</code>	mushroom-body learning curves, controls, sweeps
<code>results/phase9/</code>	capability battery across four scales, protocol fingerprints, costs
<code>results/energy/</code>	hardware measurement, biological model, per-scale curves, precision note
<code>results/assessment.json</code>	criterion verdicts, with unmeasured sub-criteria marked partial
<code>results/paper/</code>	this paper’s extracted data, figures and generated tables

Reproducing the paper.

1. extraction: `scripts/paper_collect_data.py` writes every number into `results/paper/data.json`, recording the source path for each field.
2. figures: `scripts/paper_make_figures.py` renders all nine figures, cover included.
3. tables: `scripts/paper_make_tables.py` renders the tables from the same file.
4. compile: `pdflatex paper.tex`, twice, for cross-references.

Reproducing the results themselves. Each phase script is runnable and self-documenting; run them in the order below, since later phases read earlier artifacts.

- **data and gate:** `scripts/phase0_build.py`, `scripts/phase0_gate.py`
- **geometry:** `scripts/phase4_geometry.py`
- **scaling:** `scripts/phase5_downscale.py`, `scripts/phase6_upscale.py`
- **closure:** `scripts/phase7_closure.py`
- **dynamics and learning:** `scripts/phase1_lif.py`, `scripts/phase8_learning.py`
- **capability:** `scripts/phase9_capability.py`
- **Vyb and GPU:** `scripts/vyb_des_check.py`, `scripts/phase3_gpu_check.py`
- **energy:** `scripts/phase_energy_demo.py`, `scripts/phase13_energy_curves.py`
- **verdicts:** `scripts/assess_success_criteria.py`

C Honesty log: corrections and refuted hypotheses

Table 11: Claims made during this project that turned out to be wrong, and what replaced them. Kept because the failure modes are as informative as the results.

claim	correction
Energy curves “reproduce exactly to three decimals”.	False. The two readings compared were both taken from one run’s output. A genuine repeat differs by 3.1–6.8% in joules (spike and event counts do reproduce to 0.05%). Corrected in <code>results/energy/CURVES-NOTE.md</code> and §7.2.
GPU descriptor loads are corrupted per lane.	Refuted by a purpose-built probe: a six-word descriptor is uniform on all lanes. The real defect was a store width, not a load. Recorded in <code>results/phase3/SUMMARY.txt</code> .
The last descriptor words are lost.	Refuted: the same kernel reads those words correctly when the value is stored rather than computed. Retracted in place.
A control kernel proved the store lowering was innocent.	False negative: that kernel also stored lane-dependent values, but its stores happened to be ordered so each lane’s own value landed last, masking the clobber. The PTX evidence that we had initially dismissed was correct.
Minimum success is supported.	The first version of the project assessment tool said so with the decisive $2\times$ closure test absent. The tool now counts any unmeasured sub-criterion as PARTIAL, and a missing test can never yield SUPPORTED.
Capability metrics were “constant across every measured scale”.	Only one scale had been measured, making every such verdict vacuous. Corrected in the phase-9 summary; the battery was re-run at $2\times$, $5\times$ and $10\times$.
Capability-per-watt used current capability values.	It had copied superseded protocol-v2 values; re-measured against the current protocol.
A hand-written expectation vector for the accumulate-kernel probe.	The kernel was right and our expectation was wrong (one term mis-multiplied). Fixed by recomputing the expected values from the network definition.

References

- [1] Dorkenwald, S. *et al.* Neuronal wiring diagram of an adult brain. *Nature* **634**, 124–138 (2024). doi:10.1038/s41586-024-07558-y.
- [2] Lin, A. *et al.* Network statistics of the whole-brain connectome of *Drosophila*. *Nature* **634**, 153–165 (2024). doi:10.1038/s41586-024-07968-y. (Source of the published topology values we gate against; computed on the v630 snapshot, 127,978 neurons, five-synapse threshold.)
- [3] FlyWire FAFB v783 data release: connectivity, root ids and annotations. Zenodo (2024). doi:10.5281/zenodo.10676866.
- [4] García-Pérez, G., Boguñá, M. & Serrano, M. Á. Multiscale unfolding of real networks by geometric renormalization. *Nature Physics* **14**, 583–589 (2018). doi:10.1038/s41567-018-0072-5. (The conceptual basis for the coarse-graining and subdivision operators used here.)
- [5] Panda, K. *et al.* Direct quantification of the metabolic heat output of individual *Drosophila* brains. *Cell Reports Methods* (2026). doi:10.1016/j.crmeth.2026.101501. (Biological energy anchor: ≈ 256 nW per 10-day-old female brain.)

- [6] Panda, K. *et al.* High-resolution quantification of metabolic heat output from individual live *Drosophila* brains. *bioRxiv* (2025). doi:10.1101/2025.08.08.669302.
- [7] Herculano-Houzel, S. Scaling of brain metabolism with a fixed energy budget per neuron: implications for neuronal activity, plasticity and evolution. *PLoS ONE* **6**, e17514 (2011). doi:10.1371/journal.pone.0017514. (Cross-check only; explicitly not a *Drosophila* measurement.)
- [8] Attwell, D. & Laughlin, S. B. An energy budget for signaling in the gray matter of the brain. *J. Cereb. Blood Flow Metab.* **21**, 1133–1145 (2001). doi:10.1097/00004647-200110000-00001. (Per-spike and per-vesicle cost estimates used in the biological track.)
- [9] Scheffer, L. K. *et al.* The physical design of biological systems — insights from the fly brain (2021). doi:10.1145/3439706.3446898. (Fly-brain power cross-check.)
- [10] Vyb issue #270: “with a 13-word kernel descriptor, a value computed from the later words is lost on lanes 1..7”. <https://github.com/rickenator/Vyb/issues/270>.
- [11] Vyb pull request #273, fix/270-st-intrinsic-width: kernel-mode `st_*` now uses the intrinsic’s element type. <https://github.com/rickenator/Vyb>.